

Homework 6

Iman Tabrizian
ECE1508

March 6, 2019

In this lab I learned how to implement a VNF (firewall) to protect a wordpress application. I learned that a lot of technology is required to enable simple software-based network function. I learned the basics about iptables' chains and tables and how to enable SNAT using SNAT.

What I did in this lab was implementation of a software-based firewall application. Through the means of SDN I redirected the traffic to the VM containing the VNF (firewall) and applied made the traffic to go through the snort. In this way I was able to apply simple rules to the traffic and provide better security for the Wordpress application.

Part 1

show the results of the ping tests from h1 to h2 and to h3.

```
ubuntu@netsoft17-h1:~$ ping -c 4 192.168.200.11
PING 192.168.200.11 (192.168.200.11) 56(84) bytes of data.
64 bytes from 192.168.200.11: icmp_seq=1 ttl=64 time=512 ms
64 bytes from 192.168.200.11: icmp_seq=2 ttl=64 time=4.25 ms
64 bytes from 192.168.200.11: icmp_seq=3 ttl=64 time=4.80 ms
64 bytes from 192.168.200.11: icmp_seq=4 ttl=64 time=5.05 ms

--- 192.168.200.11 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3004ms
rtt min/avg/max/mdev = 4.251/131.731/512.819/220.021 ms

ubuntu@netsoft17-h1:~$ ping -c 4 192.168.200.12
PING 192.168.200.12 (192.168.200.12) 56(84) bytes of data.
64 bytes from 192.168.200.12: icmp_seq=1 ttl=64 time=12.2 ms
64 bytes from 192.168.200.12: icmp_seq=2 ttl=64 time=4.64 ms
64 bytes from 192.168.200.12: icmp_seq=3 ttl=64 time=4.46 ms
64 bytes from 192.168.200.12: icmp_seq=4 ttl=64 time=4.71 ms

--- 192.168.200.12 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3005ms
rtt min/avg/max/mdev = 4.469/6.509/12.217/3.297 ms
```

```

ubuntu@netsoft17-h1:~$ sudo tcpdump -n -i eth0 icmp
sudo: unable to resolve host netsoft17-h1
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on eth0, link-type EN10MB (Ethernet), capture size 262144 bytes
19:04:43.185092 IP 10.12.125.12 > 172.217.1.14: ICMP echo request, id 2750, seq 1, length 64
19:04:43.191069 IP 172.217.1.14 > 10.12.125.12: ICMP echo reply, id 2750, seq 1, length 64
19:04:44.183007 IP 10.12.125.12 > 172.217.1.14: ICMP echo request, id 2750, seq 2, length 64
19:04:44.185408 IP 172.217.1.14 > 10.12.125.12: ICMP echo reply, id 2750, seq 2, length 64
19:04:45.185000 IP 10.12.125.12 > 172.217.1.14: ICMP echo request, id 2750, seq 3, length 64
19:04:45.187349 IP 172.217.1.14 > 10.12.125.12: ICMP echo reply, id 2750, seq 3, length 64
19:04:46.187649 IP 10.12.125.12 > 172.217.1.14: ICMP echo request, id 2750, seq 4, length 64
19:04:46.190520 IP 172.217.1.14 > 10.12.125.12: ICMP echo reply, id 2750, seq 4, length 64
19:04:47.188476 IP 10.12.125.12 > 172.217.1.14: ICMP echo request, id 2750, seq 5, length 64
19:04:47.191188 IP 172.217.1.14 > 10.12.125.12: ICMP echo reply, id 2750, seq 5, length 64
19:04:48.190542 IP 10.12.125.12 > 172.217.1.14: ICMP echo request, id 2750, seq 6, length 64
19:04:48.193038 IP 172.217.1.14 > 10.12.125.12: ICMP echo reply, id 2750, seq 6, length 64

```

Figure 1: TCP Dump h1

Part 2

include screenshots clearly showing the correctness of your implementation: a routing table from either h2 or h3, and the relevant iptables entry for turning h1 into a NAT router

```

ubuntu@netsoft17-h2:~$ route
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
default          192.168.200.10 0.0.0.0         UG    0      0      0 br1-internal
10.12.125.0      *               255.255.255.0   U     0      0      0 eth0
172.17.0.0       *               255.255.0.0     U     0      0      0 docker0
192.168.200.0    *               255.255.255.0   U     0      0      0 br1-internal

```

Figure 2: h2 routing table

```

ubuntu@netsoft17-h3:~$ route
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
default          192.168.200.10 0.0.0.0         UG    0      0      0 br1-internal
10.12.125.0      *               255.255.255.0   U     0      0      0 eth0
172.17.0.0       *               255.255.0.0     U     0      0      0 docker0
192.168.200.0    *               255.255.255.0   U     0      0      0 br1-internal

```

Figure 3: h3 routing table

```

sudo iptables -t nat -A POSTROUTING -o eth0 -j MASQUERADE
sudo iptables -P FORWARD ACCEPT

```

Chain POSTROUTING (policy ACCEPT)

target	prot	opt	source	destination
MASQUERADE	all	--	172.17.0.0/16	anywhere
MASQUERADE	all	--	anywhere	anywhere

```

ubuntu@netsoft17-h2:~$ route
Kernel IP routing table

```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
default	192.168.200.10	0.0.0.0	UG	0	0	0	br1-internal
10.12.125.0	*	255.255.255.0	U	0	0	0	eth0
172.17.0.0	*	255.255.0.0	U	0	0	0	docker0
192.168.200.0	*	255.255.255.0	U	0	0	0	br1-internal

```
ubuntu@netsoft17-h3:~$ route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
default	192.168.200.10	0.0.0.0	UG	0	0	0	br1-internal
10.12.125.0	*	255.255.255.0	U	0	0	0	eth0
172.17.0.0	*	255.255.0.0	U	0	0	0	docker0
192.168.200.0	*	255.255.255.0	U	0	0	0	br1-internal

Part 3.3

```
in_port={port num of vxlan iface},priority={something high},dl_dst={mac
    of internal iface created by saviOverlay},actions=output:{port num of
internal iface created by saviOverlay}
```

This rule redirects any traffic related to the internal interface to the internal internal interface.

```
in_port={port num of vxlan iface},priority={something
high},dl_dst=01:00:00:00:00:00/01:00:00:00:00:00,actions=output:{
port num of internal iface created by saviOverlay }
```

This rule redirects any broadcast traffic related to the internal interface to the internal interface.

```
in_port={port num of internal iface created by saviOverlay},actions=output:{port num of
```

This rule redirects any traffic from the internal interface to the vxlan interface to be delivered to other switches. This traffic is any traffic other than the traffic which was directed using SDN rules installed later on. The higher priority is because it is not required for other traffics to be passed inside the snort.

```
in_port={port num of vxlan iface},priority={something lower than
before},actions=output:{port num of snort ingress iface}
```

This rule redirects any traffic not related to internal interface to the snort ingress port. This is basically the traffic which was directed using SDN and is actually destined for h3. The lower priority is used because we don't want to interfere with a traffic which was not destined for h3.

```
in_port={port num of vxlan iface},priority={something lower than
before},actions=output:{port num of snort ingress iface}
```

```
in_port={port num of snort egress iface},priority={something lower than
before},actions=output:{port num of vxlan iface}
```

This rule redirects any traffic coming out of the snort ingress port to be delivered to the destination.

```
ubuntu@imant-h3:~$ sudo ovs-ofctl show br1
sudo: unable to resolve host imant-h3
OFPT_FEATURES_REPLY (xid=0x2): dpid:0000164b6e819f44
n_tables:254, n_buffers:256
capabilities: FLOW_STATS TABLE_STATS PORT_STATS QUEUE_STATS ARP_MATCH_IP
actions: output enqueue set_vlan_vid set_vlan_pcp strip_vlan mod_dl_src mod_dl_dst mod_n
1(br1-internal): addr:0a:ca:8a:da:ec:a4
    config:      0
    state:       0
    speed: 0 Mbps now, 0 Mbps max
2(imant-h3-imant-): addr:16:96:3d:4d:aa:74
    config:      0
    state:       0
    speed: 0 Mbps now, 0 Mbps max
3(snort-1): addr:72:7a:eb:72:6c:11
    config:      PORT_DOWN
    state:       LINK_DOWN
    speed: 0 Mbps now, 0 Mbps max
4(snort-2): addr:86:97:05:2d:3f:c4
    config:      PORT_DOWN
    state:       LINK_DOWN
    speed: 0 Mbps now, 0 Mbps max
LOCAL(br1): addr:16:4b:6e:81:9f:44
    config:      PORT_DOWN
    state:       LINK_DOWN
    speed: 0 Mbps now, 0 Mbps max
OFPT_GET_CONFIG_REPLY (xid=0x4): frags=normal miss_send_len=0
```

```
ubuntu@imant-h3:~$ sudo ovs-ofctl dump-flows br1
sudo: unable to resolve host imant-h3
NXST_FLOW reply (xid=0x4):
    cookie=0x0, duration=140.013s, table=0, n_packets=443, n_bytes=39886, idle_age=0, priority=
    cookie=0x0, duration=127.489s, table=0, n_packets=140, n_bytes=7140, idle_age=1, priority=
    cookie=0x0, duration=117.099s, table=0, n_packets=182, n_bytes=20356, idle_age=0, in_po
    cookie=0x0, duration=105.494s, table=0, n_packets=0, n_bytes=0, idle_age=105, priority=
    cookie=0x0, duration=94.100s, table=0, n_packets=0, n_bytes=0, idle_age=94, priority=50
    cookie=0x0, duration=20585.874s, table=0, n_packets=25253, n_bytes=1406768, idle_age=11
```

inport

Part 3.4

```
reject tcp any any -i any 80 (content:"inject"; nocase; msg:"accessed forbidden pages!!";  
sid:5000000;)
```

The first of the rule tells the action to be applied to the matching rule in this case *reject*. The second word tells which protocol does this rule apply to. In this case the protocol is *ip*. The third word tells what is the source ip address. In this case source ip address is *any*. The fourth word specifies the source port number. In this case source port number is *any*. The part after arrow specifies the destination specification. Also, the direction of the arrow indicates the direction of interest. The first word after arrow is destination ip address which in this case is *any*. The second word after arrow is the destination port number which in this case is *80*. The part in the paranthesis specifies the specific options regarding the rule. The *nocase* option is used to deactivate any case sensitivity in the content rule. The *msg* specifies the message to be printed along the packet dump. The *sid* is used for specifying a particular snort rule. This can be used by external plugins to identify snort rules.

Part 4.4

You can find the code for installing flows in the appendix section. Unfortunately, because my VMs were deleted (3 times), I didn't find enough time to go through the lab again from the beginning. This is the reason why flow dumps are missing from this section.

```
ubuntu@imant-sw3:~$ sudo ovs-ofctl dump-flows br1  
sudo: unable to resolve host imant-sw3  
NXST_FLOW reply (xid=0x4):  
  cookie=0x0, duration=21387.397s, table=0, n_packets=23729, n_bytes=1210179, idle_age=1,  
  cookie=0x0, duration=79.049s, table=0, n_packets=30, n_bytes=6274, idle_age=45, priorit  
  cookie=0x0, duration=24.857s, table=0, n_packets=0, n_bytes=0, idle_timeout=60, hard_ti  
  cookie=0x0, duration=24.803s, table=0, n_packets=0, n_bytes=0, idle_timeout=60, hard_ti  
  
sudo: unable to resolve host imant-sw1  
NXST_FLOW reply (xid=0x4):  
  cookie=0x0, duration=21433.695s, table=0, n_packets=47572, n_bytes=2426172, idle_age=0,  
  cookie=0x0, duration=3249.019s, table=0, n_packets=0, n_bytes=0, idle_age=3249, priorit  
  cookie=0x0, duration=3249.018s, table=0, n_packets=0, n_bytes=0, idle_age=3249, priorit  
  cookie=0x0, duration=130.100s, table=0, n_packets=30, n_bytes=6274, idle_age=96, priori  
  cookie=0x0, duration=130.099s, table=0, n_packets=30, n_bytes=3829, idle_age=96, priori
```

Part 4.5

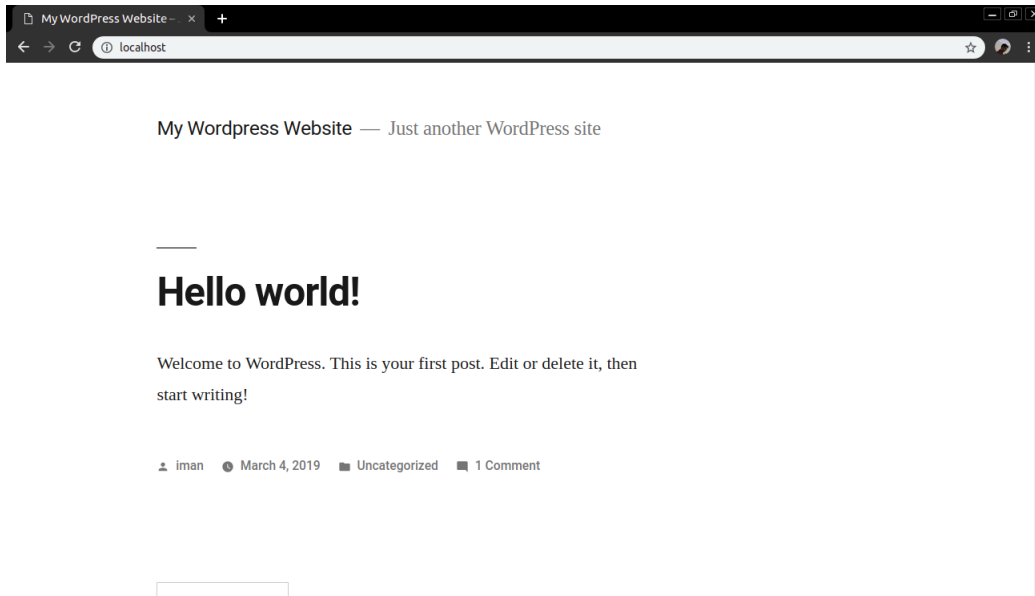


Figure 4: Loading WordPress

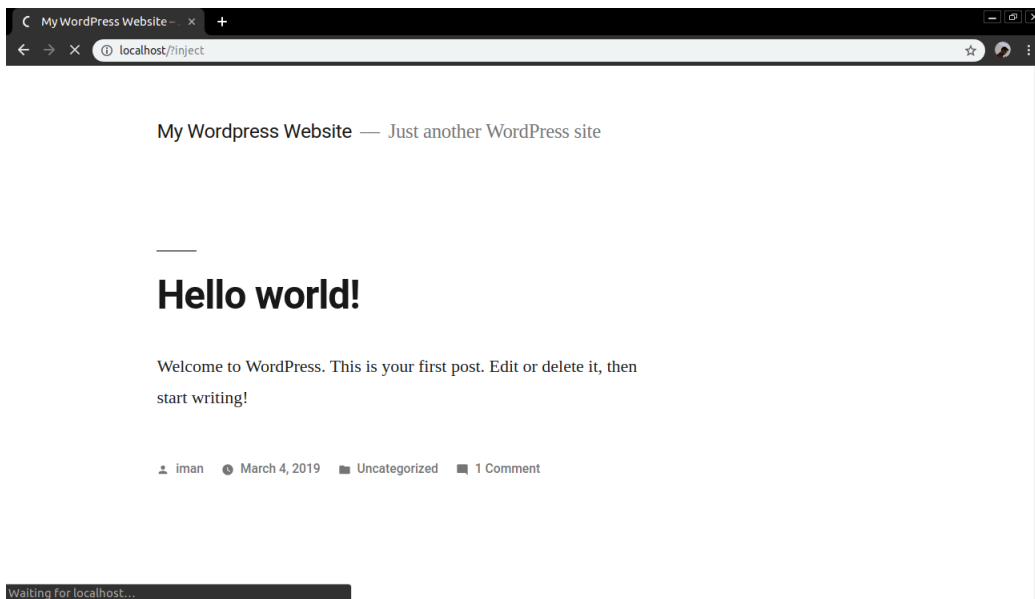


Figure 5: Not Loading WordPress if we use inject

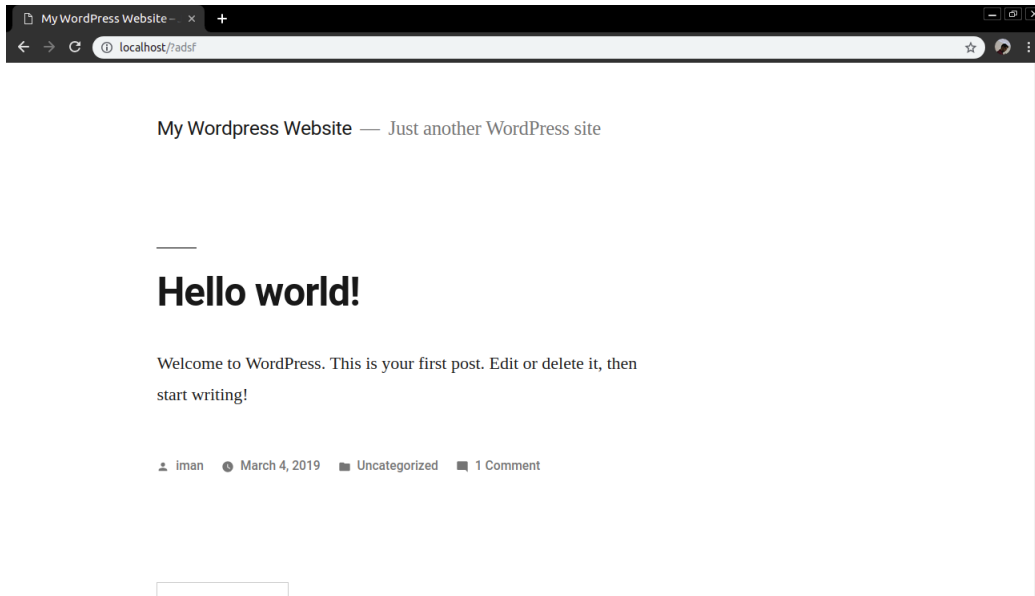


Figure 6: Other endpoints working




Figure 7: Snort logs showing blocking of connection

Appendix

```
import ryu_ofctl

h2 = "5a:b5:e9:15:cc:ec"
h3 = "0a:ca:8a:da:ec:a4"
h1 = "76:eb:89:2e:25:70"

sw1 = '00002a38cb041745'
sw2 = "0000e601655db64f"
sw3 = "0000fa9ecd0fb541"

sw1_flows = [
    ryu_ofctl.FlowEntry(),
    ryu_ofctl.FlowEntry(),
    ryu_ofctl.FlowEntry()
]

# Flow 1 (h1 -> h3 for initial inspection) (sw1)
sw1_flows[0].dl_src = h1
sw1_flows[0].dl_dst = h2
sw1_flows[0].in_port = 1
sw1_flows[0].dl_type = 0x800
sw1_flows[0].nw_proto = 0x6
sw1_flows[0].tp_dst = 80
sw1_flows[0].priority = 60000
sw1_flows[0].addAction(ryu_ofctl.OutputAction(2))
ryu_ofctl.insertFlow(sw1, sw1_flows[0])

# Flow 2 (h3 -> h2 for passing after processing) (sw1)
sw1_flows[1].dl_src = h1
sw1_flows[1].dl_dst = h2
sw1_flows[1].in_port = 2
sw1_flows[1].dl_type = 0x800
sw1_flows[1].nw_proto = 0x6
sw1_flows[1].tp_dst = 80
sw1_flows[1].priority = 60000
sw1_flows[1].addAction(ryu_ofctl.OutputAction(3))
ryu_ofctl.insertFlow(sw1, sw1_flows[1])

# Flow 3 (h3 -> h2 for passing after processing) (sw3)
sw1_flows[2].dl_src = h1
sw1_flows[2].dl_dst = h2
sw1_flows[2].in_port = 1
```



```
sw1_flows[2].dl_type = 0x800
sw1_flows[2].nw_proto = 0x6
sw1_flows[2].tp_dst = 80
sw1_flows[2].priority = 60000
sw1_flows[2].addAction(ryu_ofctl.OutputAction(2))
ryu_ofctl.insertFlow(sw3, sw1_flows[2])
```